

# 301107 - Analytics Programming

Nick Tothill

*Practical class 1 - Excel; setting up R; R environment; basic data handling*

## Task 1. Clean up the Iris data

Use Excel to produce basic summaries of the iris data (given in vUWS, under Learning Materials; week02; practical). This famous (Fisher's or Anderson's) iris data set gives the measurements in centimetres of the variables sepal length and width and petal length and width, respectively, for 50 flowers from each of 3 species of iris. The species are *Iris setosa*, *versicolor*, and *virginica*.

The data look like:

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

Hover, during the data collection process, something went wrong. As a result, some values are missing, and some values are incorrect (with values  $\leq 0$ ).

This is a common problem in data science. So the first thing to do is to

- **clean up the data** by removing any record with **either missing values or incorrect values**
- copy the records with *good* values to another sheet and carry on the next task

## Task 2. Summarise the Iris data

Given now you have cleaned the data, try to work out some basic statistics of the flowers summarised according to different species. So you produce a table that looks like the following (can be in different format):

```
[1] "Species: setosa"
  Sepal#Length  Sepal#Width  Petal#Length  Petal#Width
Min.   :4.300    Min.   :2.300   Min.   :1.000   Min.   :0.1000
1st Qu.:4.750    1st Qu.:3.100   1st Qu.:1.350   1st Qu.:0.2000
Median :5.000    Median :3.400   Median :1.400   Median :0.2000
Mean   :4.957    Mean   :3.346   Mean   :1.437   Mean   :0.2371
3rd Qu.:5.100    3rd Qu.:3.600   3rd Qu.:1.500   3rd Qu.:0.3000
Max.   :5.800    Max.   :4.000   Max.   :1.900   Max.   :0.5000

[1] "Species: versicolor"
  Sepal#Length  Sepal#Width  Petal#Length  Petal#Width
Min.   :4.900    Min.   :2.000   Min.   :3.000   Min.   :1.000
1st Qu.:5.500    1st Qu.:2.500   1st Qu.:3.900   1st Qu.:1.200
Median :5.800    Median :2.800   Median :4.200   Median :1.300
Mean   :5.885    Mean   :2.741   Mean   :4.171   Mean   :1.295
3rd Qu.:6.200    3rd Qu.:3.000   3rd Qu.:4.500   3rd Qu.:1.400
Max.   :7.000    Max.   :3.200   Max.   :4.900   Max.   :1.800


[1] "Species: virginica"
```

Sepal#Length	Sepal#Width	Petal#Length	Petal#Width
Min. :4.900	Min. :2.20	Min. :4.500	Min. :1.400
1st Qu.:6.175	1st Qu.:2.80	1st Qu.:5.100	1st Qu.:1.800
Median :6.450	Median :3.00	Median :5.600	Median :2.000
Mean :6.590	Mean :2.98	Mean :5.575	Mean :2.002
3rd Qu.:7.125	3rd Qu.:3.20	3rd Qu.:5.825	3rd Qu.:2.225
Max. :7.900	Max. :3.80	Max. :6.900	Max. :2.500

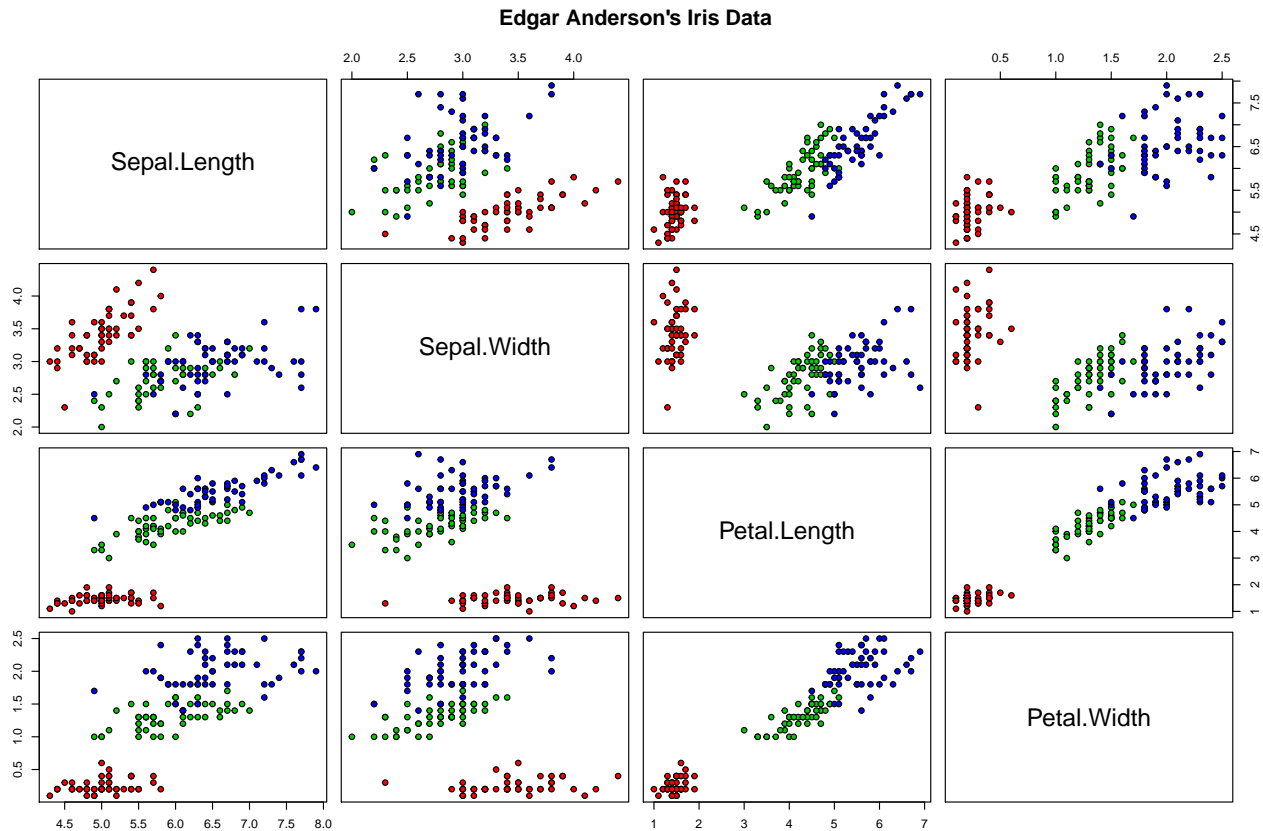
You may notice there are 6 statistics shown above. In this practical class, you *only need to find minimum, maximum and mean* for each variable.

### Task 3. Visualise the Data

Plot a figure. Use Excel's X Y (Scatter) chart to produce a plot similar to the following



However, it would be better to separate the species:



## Task 4. Starting R

Start Rstudio on your lab computer. If using your own computer, you should have Rstudio installed. If not, download and install while you're working on the practical on a lab computer.

In RStudio, try to run the following code

```
> 1+1
```

```
[1] 2
```

```
> sqrt(2)
```

```
[1] 1.414214
```

```
> print('Hello, world!')
```

```
[1] "Hello, world!"
```

## Task 5. Simple R session

Type in the following commands:

```
> x <- 1
> y <- 2
> z <- x+y
> ls()
```

```
> rm(x)
> ls()
> help(ls)
> help.start()
```

Look for *An Introduction to R*.

Go back to RStudio *R Console* panel and continue

```
> ls
> ls <- 1
> ls
> ls()
```

What happens with these different ways of writing `ls`? Notice the changes?

Continue

```
> getwd()
> save.image(file='test.RData')
> history()
```

What is `save.image(file='test.RData')` doing there?

What if I want to save all commands I used so far? (Hint: look up things related to `history` )

```
> rm(list=ls())
> load(file='test.RData')
> ls()
```

We cleared the current workspace by using `rm(list=ls())`. See what the argument `list` of function `rm()` means. What is `load()` doing?

If you have worked out how to save all commands you used in this session, you know how to save your work, which is likely to be important!

## Task 6. Practice with vectors

Try out the following code for vectors, and try to figure out what every line is doing

```
x <- seq(0,200,5)
print(paste('x is a vector of length',length(x)))
x
plot(1:length(x),x,main="Plot of vector x", xlab='Index', ylab='Values of elements in x',col=1)

print(paste('The 10th element in x is',x[10]))
cat('The first 5 elements in x are',x[1:5])
cat('The 5th, 7th, and 10th elements in x are',x[c(5,7,10)])
x1 <- x[1:4]
x2 <- x[(length(x)-3):length(x)]
x1
x1[-1]
x2
x1 + x2
-x1
x1 * x2
any(x1>5)
```

```

which(x1>5)
x1[which(x1>5)]
x1[1] <- NA
x1
which(is.na(x1))
x1[-which(is.na(x1))]
all(x2<200)

x1 <- x1[-1]
x1
x1-5
x1*-1
sum(x1)
rep(1,length(x2))

```

There are quite a few indexing and manipulating functions in the above code, as well as a plotting function `plot()`. Find out what they do to a vector (hint: use `help()`.)

After you understand the above, finish the following task, using the functions you have learned to make the code as simple as possible.

1. generate a random vector of length 10 (the function you need is `runif(10)`);
2. calculate the sum of the samples stored in the vector;
3. calculate the mean of the samples;
4. find all samples that are no less than 0.5 and calculate their mean;
5. find all samples that are less than 0.5 and calculate their mean.

You should have the results similar to the following:

My vector is

```

0.0912563 0.7230549 0.7220213 0.8507309 0.48788 0.7993466
0.2850184 0.0008180665 0.128345 0.5040901

```

Sum of samples: 4.592562


Mean of samples: 0.4592562

Mean of samples no less than 0.5: 0.7198488

Mean of samples less than 0.5: 0.1986636

your vector will have different elements in it, as, every time you run `runif` it will generate different values drawn from a uniform distribution between 0 and 1. So your results will be different from the above.

You can plot the vector here to get a visualisation of the vector like the following. The command to use is `plot`. The line in the middle (which is the mean of the data in the vector) is produced by another command `abline`. Use the help pages in R to see if you can figure out how to make a plot like the one shown below.



## Task 7. R self learning package swirl

A. install R package `swirl`

1. Start RStudio

2. Type in `install.packages('swirl')`. It will download and install the package.

3. load in the library by using `library(swirl)`

B. Have a `swirl` learning session

**N.B.** The learning module may take a long time to download.